

PySpark Crash Course

January 27, 2021

```
[5]: # Load Pyspark Pkgs
import pyspark
```

```
[7]: # Methods/Attrib
dir(pyspark)
```

```
[7]: ['Accumulator',
      'AccumulatorParam',
      'BarrierTaskContext',
      'BarrierTaskInfo',
      'BasicProfiler',
      'Broadcast',
      'HiveContext',
      'MarshalSerializer',
      'PickleSerializer',
      'Profiler',
      'RDD',
      'RDDBarrier',
      'Row',
      'SQLContext',
      'SparkConf',
      'SparkContext',
      'SparkFiles',
      'SparkJobInfo',
      'SparkStageInfo',
      'StatusTracker',
      'StorageLevel',
      'TaskContext',
      '_NoValue',
      '__all__',
      '__builtins__',
      '__cached__',
      '__doc__',
      '__file__',
      '__loader__',
      '__name__',
      '__package__']
```

```

'__path__',
'__spec__',
'__version__',
'_globals',
'accumulators',
'broadcast',
'cloudpickle',
'conf',
'context',
'copy_func',
'files',
'find_spark_home',
'heapq3',
'java_gateway',
'join',
'keyword_only',
'profiler',
'rdd',
'rddsampler',
'resource',
'resultiterable',
'serializers',
'shuffle',
'since',
'sql',
'statcounter',
'status',
'storagelevel',
'taskcontext',
'traceback_utils',
'types',
'util',
'version',
'wraps']

```

Working with DataFrames in PySpark

- Read DataSet(CSV)
- Create DataFrame

Tips

- SparkSession
- SparkContext :sc
- SqlContext

```

[8]: # Create A SparkSession
from pyspark.sql import SparkSession

```

```
spark = SparkSession.builder.appName("PySparkTut").getOrCreate()
```

```
[9]: !ls
```

```
'PySpark Crash Course.ipynb'    work
```

```
[10]: !wget https://raw.githubusercontent.com/Jcharis/common_ml_datasets_explorer_app/
      ↪master/datasets/diamonds.csv
```

```
--2021-01-27 21:01:43-- https://raw.githubusercontent.com/Jcharis/common_ml_dat
assets_explorer_app/master/datasets/diamonds.csv
Resolving raw.githubusercontent.com (raw.githubusercontent.com)...
151.101.112.133
Connecting to raw.githubusercontent.com
(raw.githubusercontent.com)|151.101.112.133|:443... connected.
HTTP request sent, awaiting response... 200 OK
Length: 2772143 (2.6M) [text/plain]
Saving to: 'diamonds.csv'
```

```
diamonds.csv          100%[=====>]    2.64M  3.81MB/s    in 0.7s
```

```
2021-01-27 21:01:44 (3.81 MB/s) - 'diamonds.csv' saved [2772143/2772143]
```

```
[11]: !ls
```

```
diamonds.csv  'PySpark Crash Course.ipynb'    work
```

```
[12]: # Read A DataSet without header
df = spark.read.csv('diamonds.csv')
```

```
[13]: # Preview dataset
df.show()
```

```
+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
|_c0|      _c1|_c2|      _c3|_c4|_c5|_c6|_c7|_c8|_c9|
+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
|carat|      cut|color|clarity|depth|table|price|  x|  y|  z|
| 0.23|   Ideal|  E|   SI2| 61.5|  55| 326|3.95|3.98|2.43|
| 0.21| Premium|  E|  SI1| 59.8|  61| 326|3.89|3.84|2.31|
| 0.23|    Good|  E|  VS1| 56.9|  65| 327|4.05|4.07|2.31|
| 0.29| Premium|  I|  VS2| 62.4|  58| 334| 4.2|4.23|2.63|
| 0.31|    Good|  J|  SI2| 63.3|  58| 335|4.34|4.35|2.75|
| 0.24|Very Good|  J| VVS2| 62.8|  57| 336|3.94|3.96|2.48|
| 0.24|Very Good|  I| VVS1| 62.3|  57| 336|3.95|3.98|2.47|
| 0.26|Very Good|  H|  SI1| 61.9|  55| 337|4.07|4.11|2.53|
| 0.22|    Fair|  E|  VS2| 65.1|  61| 337|3.87|3.78|2.49|
| 0.23|Very Good|  H|  VS1| 59.4|  61| 338| 4|4.05|2.39|
```

0.3	Good	J	SI1	64	55	339	4.25	4.28	2.73
0.23	Ideal	J	VS1	62.8	56	340	3.93	3.9	2.46
0.22	Premium	F	SI1	60.4	61	342	3.88	3.84	2.33
0.31	Ideal	J	SI2	62.2	54	344	4.35	4.37	2.71
0.2	Premium	E	SI2	60.2	62	345	3.79	3.75	2.27
0.32	Premium	E	I1	60.9	58	345	4.38	4.42	2.68
0.3	Ideal	I	SI2	62	54	348	4.31	4.34	2.68
0.3	Good	J	SI1	63.4	54	351	4.23	4.29	2.7
0.3	Good	J	SI1	63.8	56	351	4.23	4.26	2.71

only showing top 20 rows

```
[62]: # Read A DataSet with header/column names
df = spark.read.csv('diamonds.csv',header=True)
```

```
[15]: df.show()
```

carat	cut	color	clarity	depth	table	price	x	y	z
0.23	Ideal	E	SI2	61.5	55	326	3.95	3.98	2.43
0.21	Premium	E	SI1	59.8	61	326	3.89	3.84	2.31
0.23	Good	E	VS1	56.9	65	327	4.05	4.07	2.31
0.29	Premium	I	VS2	62.4	58	334	4.2	4.23	2.63
0.31	Good	J	SI2	63.3	58	335	4.34	4.35	2.75
0.24	Very Good	J	VVS2	62.8	57	336	3.94	3.96	2.48
0.24	Very Good	I	VVS1	62.3	57	336	3.95	3.98	2.47
0.26	Very Good	H	SI1	61.9	55	337	4.07	4.11	2.53
0.22	Fair	E	VS2	65.1	61	337	3.87	3.78	2.49
0.23	Very Good	H	VS1	59.4	61	338	4	4.05	2.39
0.3	Good	J	SI1	64	55	339	4.25	4.28	2.73
0.23	Ideal	J	VS1	62.8	56	340	3.93	3.9	2.46
0.22	Premium	F	SI1	60.4	61	342	3.88	3.84	2.33
0.31	Ideal	J	SI2	62.2	54	344	4.35	4.37	2.71
0.2	Premium	E	SI2	60.2	62	345	3.79	3.75	2.27
0.32	Premium	E	I1	60.9	58	345	4.38	4.42	2.68
0.3	Ideal	I	SI2	62	54	348	4.31	4.34	2.68
0.3	Good	J	SI1	63.4	54	351	4.23	4.29	2.7
0.3	Good	J	SI1	63.8	56	351	4.23	4.26	2.71
0.3	Very Good	J	SI1	62.7	59	351	4.21	4.27	2.66

only showing top 20 rows

```
[16]: # Columns
df.columns
```

```
[16]: ['carat', 'cut', 'color', 'clarity', 'depth', 'table', 'price', 'x', 'y', 'z']
```

```
[20]: # Shape (rows + columns)
(df.count() ,len(df.columns))
```

```
[20]: (53940, 10)
```

```
[21]: # Number of columns
len(df.columns)
```

```
[21]: 10
```

```
[22]: # Number of rows
df.count()
```

```
[22]: 53940
```

```
[24]: # Descriptive Analysis
df.describe().show()
```

```
+-----+-----+-----+-----+-----+-----+
-----+-----+-----+-----+-----+-----+
----+
|summary|          carat|          cut|color|clarity|          depth|
table|          price|          x|          y|
z|
+-----+-----+-----+-----+-----+-----+
-----+-----+-----+-----+-----+-----+
----+
| count|          53940|          53940|53940|          53940|          53940|
53940|          53940|          53940|          53940|
53940|
| mean|0.7979397478679852|          null| null|          null| 61.74940489432624|
57.45718390804603|3932.799721913237| 5.731157211716609|
5.734525954764462|3.5387337782723316|
| stddev|0.4740112444054196|          null| null|          null|1.4326213188336525|2.2344905
628213247|3989.439738146397|1.1217607467924915|1.1421346741235616|0.705698846949
9883|
| min|          0.2|          Fair|          D|          I1|          43|
43|          1000|          0|          0|          0|          0|
| max|          5.01|Very Good|          J|          VVS2|          79|
95|          9999|          9.86|          9.94|          8.06|
+-----+-----+-----+-----+-----+-----+
-----+-----+-----+-----+-----+-----+
----+
```

```
[25]: # Pick a column & Get summary/describe a selected column
df.describe('carat').show()
```

```
+-----+-----+
|summary|      carat|
+-----+-----+
|  count|      53940|
|   mean|0.7979397478679852|
| stddev|0.4740112444054196|
|    min|         0.2|
|    max|         5.01|
+-----+-----+
```

```
[26]: # Preview the First Row
df.first()
```

```
[26]: Row(carat='0.23', cut='Ideal', color='E', clarity='SI2', depth='61.5',
table='55', price='326', x='3.95', y='3.98', z='2.43')
```

```
[31]: # Preview the first 10 rows
# Like a list
df.head(10)
```

```
[31]: [Row(carat='0.23', cut='Ideal', color='E', clarity='SI2', depth='61.5',
table='55', price='326', x='3.95', y='3.98', z='2.43'),
Row(carat='0.21', cut='Premium', color='E', clarity='SI1', depth='59.8',
table='61', price='326', x='3.89', y='3.84', z='2.31'),
Row(carat='0.23', cut='Good', color='E', clarity='VS1', depth='56.9',
table='65', price='327', x='4.05', y='4.07', z='2.31'),
Row(carat='0.29', cut='Premium', color='I', clarity='VS2', depth='62.4',
table='58', price='334', x='4.2', y='4.23', z='2.63'),
Row(carat='0.31', cut='Good', color='J', clarity='SI2', depth='63.3',
table='58', price='335', x='4.34', y='4.35', z='2.75'),
Row(carat='0.24', cut='Very Good', color='J', clarity='VVS2', depth='62.8',
table='57', price='336', x='3.94', y='3.96', z='2.48'),
Row(carat='0.24', cut='Very Good', color='I', clarity='VVS1', depth='62.3',
table='57', price='336', x='3.95', y='3.98', z='2.47'),
Row(carat='0.26', cut='Very Good', color='H', clarity='SI1', depth='61.9',
table='55', price='337', x='4.07', y='4.11', z='2.53'),
Row(carat='0.22', cut='Fair', color='E', clarity='VS2', depth='65.1',
table='61', price='337', x='3.87', y='3.78', z='2.49'),
Row(carat='0.23', cut='Very Good', color='H', clarity='VS1', depth='59.4',
table='61', price='338', x='4', y='4.05', z='2.39')]
```

```
[32]: # Method 2: Useful Action with show()
# Show first 10 datapoints
```

```
df.show(10)
```

```
+-----+-----+-----+-----+-----+-----+-----+-----+-----+
|carat|      cut|color|clarity|depth|table|price|  x|  y|  z|
+-----+-----+-----+-----+-----+-----+-----+-----+-----+
| 0.23|   Ideal|  E|   SI2| 61.5|  55| 326|3.95|3.98|2.43|
| 0.21|  Premium|  E|   SI1| 59.8|  61| 326|3.89|3.84|2.31|
| 0.23|    Good|  E|   VS1| 56.9|  65| 327|4.05|4.07|2.31|
| 0.29|  Premium|  I|   VS2| 62.4|  58| 334| 4.2|4.23|2.63|
| 0.31|    Good|  J|   SI2| 63.3|  58| 335|4.34|4.35|2.75|
| 0.24|Very Good|  J|  VVS2| 62.8|  57| 336|3.94|3.96|2.48|
| 0.24|Very Good|  I|  VVS1| 62.3|  57| 336|3.95|3.98|2.47|
| 0.26|Very Good|  H|   SI1| 61.9|  55| 337|4.07|4.11|2.53|
| 0.22|    Fair|  E|   VS2| 65.1|  61| 337|3.87|3.78|2.49|
| 0.23|Very Good|  H|   VS1| 59.4|  61| 338| 4|4.05|2.39|
+-----+-----+-----+-----+-----+-----+-----+-----+-----+
only showing top 10 rows
```

```
[33]: # Get Last Rows
df.tail(5)
```

```
[33]: [Row(carat='0.72', cut='Ideal', color='D', clarity='SI1', depth='60.8',
table='57', price='2757', x='5.75', y='5.76', z='3.5'),
Row(carat='0.72', cut='Good', color='D', clarity='SI1', depth='63.1',
table='55', price='2757', x='5.69', y='5.75', z='3.61'),
Row(carat='0.7', cut='Very Good', color='D', clarity='SI1', depth='62.8',
table='60', price='2757', x='5.66', y='5.68', z='3.56'),
Row(carat='0.86', cut='Premium', color='H', clarity='SI2', depth='61',
table='58', price='2757', x='6.15', y='6.12', z='3.74'),
Row(carat='0.75', cut='Ideal', color='D', clarity='SI2', depth='62.2',
table='55', price='2757', x='5.83', y='5.87', z='3.64')]
```

0.0.1 Selection of columns

- .select ##### Note
- Dot & Bracket Notation only gives the column name not the entire column
 - ['colA']*
 - .colA*

```
[35]: # List all Columns
df.columns
```

```
[35]: ['carat', 'cut', 'color', 'clarity', 'depth', 'table', 'price', 'x', 'y', 'z']
```

```
[37]: # Select A Column
df.select('carat').show()
```

```

+-----+
|carat|
+-----+
| 0.23|
| 0.21|
| 0.23|
| 0.29|
| 0.31|
| 0.24|
| 0.24|
| 0.26|
| 0.22|
| 0.23|
| 0.3|
| 0.23|
| 0.22|
| 0.31|
| 0.2|
| 0.32|
| 0.3|
| 0.3|
| 0.3|
| 0.3|
+-----+

```

only showing top 20 rows

```

[40]: # Select A Column irrespective of column word case
      # will work irrespective of the case of the column once it is found within the
      ↪ dataset
      df.select('CARAT').show()

```

```

+-----+
|CARAT|
+-----+
| 0.23|
| 0.21|
| 0.23|
| 0.29|
| 0.31|
| 0.24|
| 0.24|
| 0.26|
| 0.22|
| 0.23|
| 0.3|
| 0.23|
| 0.22|

```



```
| 0.31|
| 0.2|
| 0.32|
| 0.3|
| 0.3|
| 0.3|
| 0.3|
+-----+
```

only showing top 20 rows

```
[41]: # This is not as we would expect in pandas
# For Bracket Notation : pick column name not the entire column
df['carat']
```

```
[41]: Column<b'carat'>
```

```
[44]: # This is not as we would expect in pandas
# For Dot Notation : pick column name not the entire column
df.carat
```

```
[44]: Column<b'carat'>
```

```
[45]: # Select Multiple Columns
df.select('carat','cut').show(5)
```

```
+-----+-----+
|carat|    cut|
+-----+-----+
| 0.23| Ideal|
| 0.21|Premium|
| 0.23|  Good|
| 0.29|Premium|
| 0.31|  Good|
+-----+-----+
```

only showing top 5 rows

0.0.2 Column Filtering and Applying Conditions

- .filter
- .where

```
[46]: # Filter of Columns
# Apply A Condition
df.show(10)
```

```
+-----+-----+-----+-----+-----+-----+-----+-----+
|carat|    cut|color|clarity|depth|table|price|  x|  y|  z|
```

```

+-----+-----+-----+-----+-----+-----+-----+-----+-----+
| 0.23|    Ideal|    E|    SI2| 61.5|   55|  326|3.95|3.98|2.43|
| 0.21|   Premium|    E|    SI1| 59.8|   61|  326|3.89|3.84|2.31|
| 0.23|     Good|    E|    VS1| 56.9|   65|  327|4.05|4.07|2.31|
| 0.29|   Premium|    I|    VS2| 62.4|   58|  334| 4.2|4.23|2.63|
| 0.31|     Good|    J|    SI2| 63.3|   58|  335|4.34|4.35|2.75|
| 0.24|Very Good|    J|   VVS2| 62.8|   57|  336|3.94|3.96|2.48|
| 0.24|Very Good|    I|   VVS1| 62.3|   57|  336|3.95|3.98|2.47|
| 0.26|Very Good|    H|    SI1| 61.9|   55|  337|4.07|4.11|2.53|
| 0.22|     Fair|    E|    VS2| 65.1|   61|  337|3.87|3.78|2.49|
| 0.23|Very Good|    H|    VS1| 59.4|   61|  338| 4|4.05|2.39|
+-----+-----+-----+-----+-----+-----+-----+-----+

```

only showing top 10 rows

```
[47]: # Method 1:using filter
df.filter(df['cut'] == "Good").show()
```

```

+-----+-----+-----+-----+-----+-----+-----+-----+-----+
|carat| cut|color|clarity|depth|table|price|  x|  y|  z|
+-----+-----+-----+-----+-----+-----+-----+-----+-----+
| 0.23|Good|    E|    VS1| 56.9|   65|  327|4.05|4.07|2.31|
| 0.31|Good|    J|    SI2| 63.3|   58|  335|4.34|4.35|2.75|
| 0.3|Good|    J|    SI1| 64|   55|  339|4.25|4.28|2.73|
| 0.3|Good|    J|    SI1| 63.4|   54|  351|4.23|4.29| 2.7|
| 0.3|Good|    J|    SI1| 63.8|   56|  351|4.23|4.26|2.71|
| 0.3|Good|    I|    SI2| 63.3|   56|  351|4.26| 4.3|2.71|
| 0.23|Good|    F|    VS1| 58.2|   59|  402|4.06|4.08|2.37|
| 0.23|Good|    E|    VS1| 64.1|   59|  402|3.83|3.85|2.46|
| 0.31|Good|    H|    SI1| 64|   54|  402|4.29|4.31|2.75|
| 0.26|Good|    D|    VS2| 65.2|   56|  403|3.99|4.02|2.61|
| 0.26|Good|    D|    VS1| 58.4|   63|  403|4.19|4.24|2.46|
| 0.32|Good|    H|    SI2| 63.1|   56|  403|4.34|4.37|2.75|
| 0.32|Good|    H|    SI2| 63.8|   56|  403|4.36|4.38|2.79|
| 0.3|Good|    I|    SI1| 63.2|   55|  405|4.25|4.29| 2.7|
| 0.3|Good|    H|    SI1| 63.7|   57|  554|4.28|4.26|2.72|
| 0.26|Good|    E|   VVS1| 57.9|   60|  554|4.22|4.25|2.45|
| 0.7|Good|    E|    VS2| 57.5|   58| 2759|5.85| 5.9|3.38|
| 0.7|Good|    F|    VS1| 59.4|   62| 2759|5.71|5.76| 3.4|
| 0.7|Good|    H|   VVS2| 62.1|   64| 2767|5.62|5.65| 3.5|
| 0.71|Good|    E|    VS2| 59.2|   61| 2772| 5.8|5.88|3.46|
+-----+-----+-----+-----+-----+-----+-----+-----+

```

only showing top 20 rows

```
[48]: # Method 1:using filter
df.filter(df.carat >= 0.7).show()
```

carat	cut	color	clarity	depth	table	price	x	y	z
0.7	Ideal	E	SI1	62.5	57	2757	5.7	5.72	3.57
0.86	Fair	E	SI2	55.1	69	2757	6.45	6.33	3.52
0.7	Ideal	G	VS2	61.6	56	2757	5.7	5.67	3.5
0.71	Very Good	E	VS2	62.4	57	2759	5.68	5.73	3.56
0.78	Very Good	G	SI2	63.8	56	2759	5.81	5.85	3.72
0.7	Good	E	VS2	57.5	58	2759	5.85	5.9	3.38
0.7	Good	F	VS1	59.4	62	2759	5.71	5.76	3.4
0.96	Fair	F	SI2	66.3	62	2759	6.27	5.95	4.07
0.73	Very Good	E	SI1	61.6	59	2760	5.77	5.78	3.56
0.8	Premium	H	SI1	61.5	58	2760	5.97	5.93	3.66
0.75	Very Good	D	SI1	63.2	56	2760	5.8	5.75	3.65
0.75	Premium	E	SI1	59.9	54	2760	6	5.96	3.58
0.74	Ideal	G	SI1	61.6	55	2760	5.8	5.85	3.59
0.75	Premium	G	VS2	61.7	58	2760	5.85	5.79	3.59
0.8	Ideal	I	VS1	62.9	56	2760	5.94	5.87	3.72
0.75	Ideal	G	SI1	62.2	55	2760	5.87	5.8	3.63
0.8	Premium	G	SI1	63	59	2760	5.9	5.81	3.69
0.74	Ideal	I	VVS2	62.3	55	2761	5.77	5.81	3.61
0.81	Ideal	F	SI2	58.8	57	2761	6.14	6.11	3.6
0.8	Ideal	F	SI2	61.4	57	2761	5.96	6	3.67

only showing top 20 rows

```
[49]: # Method 2: where
df.where(df['cut'] == 'Good').show()
```

carat	cut	color	clarity	depth	table	price	x	y	z
0.23	Good	E	VS1	56.9	65	327	4.05	4.07	2.31
0.31	Good	J	SI2	63.3	58	335	4.34	4.35	2.75
0.3	Good	J	SI1	64	55	339	4.25	4.28	2.73
0.3	Good	J	SI1	63.4	54	351	4.23	4.29	2.7
0.3	Good	J	SI1	63.8	56	351	4.23	4.26	2.71
0.3	Good	I	SI2	63.3	56	351	4.26	4.3	2.71
0.23	Good	F	VS1	58.2	59	402	4.06	4.08	2.37
0.23	Good	E	VS1	64.1	59	402	3.83	3.85	2.46
0.31	Good	H	SI1	64	54	402	4.29	4.31	2.75
0.26	Good	D	VS2	65.2	56	403	3.99	4.02	2.61
0.26	Good	D	VS1	58.4	63	403	4.19	4.24	2.46
0.32	Good	H	SI2	63.1	56	403	4.34	4.37	2.75
0.32	Good	H	SI2	63.8	56	403	4.36	4.38	2.79
0.3	Good	I	SI1	63.2	55	405	4.25	4.29	2.7
0.3	Good	H	SI1	63.7	57	554	4.28	4.26	2.72

0.26	Good	E	VVS1	57.9	60	554	4.22	4.25	2.45
0.7	Good	E	VS2	57.5	58	2759	5.85	5.9	3.38
0.7	Good	F	VS1	59.4	62	2759	5.71	5.76	3.4
0.7	Good	H	VVS2	62.1	64	2767	5.62	5.65	3.5
0.71	Good	E	VS2	59.2	61	2772	5.8	5.88	3.46

+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+

only showing top 20 rows

```
[51]: # Method 2: where
# select certain columns
df.where(df['cut'] == 'Good').select('price','cut').show()
```

+-----+-----+

price	cut
-------	-----

+-----+-----+

327	Good
335	Good
339	Good
351	Good
351	Good
351	Good
402	Good
402	Good
402	Good
403	Good
403	Good
403	Good
403	Good
405	Good
554	Good
554	Good
2759	Good
2759	Good
2767	Good
2772	Good

+-----+-----+

only showing top 20 rows

```
[52]: # Unique Values
# df['cut'].unique()
df.select("cut").distinct().show()
```

+-----+-----+

cut

+-----+-----+

Premium

```
|    Ideal|
|    Good|
|    Fair|
|Very Good|
+-----+
```

0.0.3 How to Add Columns & Delete/Drop Columns

- .withColumn()
- .drop()

```
[54]: # Add Columns
df.withColumn("carat10x",df['carat'] * 10).show()
```

```
+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
--+
|carat|      cut|color|clarity|depth|table|price|  x|  y|  z|
carat10x|
+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
--+
| 0.23|    Ideal|    E|   SI2| 61.5|  55|
326|3.95|3.98|2.43|2.3000000000000003|
| 0.21|  Premium|    E|   SI1| 59.8|  61| 326|3.89|3.84|2.31|
2.1|
| 0.23|    Good|    E|   VS1| 56.9|  65|
327|4.05|4.07|2.31|2.3000000000000003|
| 0.29|  Premium|    I|   VS2| 62.4|  58| 334| 4.2|4.23|2.63|
2.9|
| 0.31|    Good|    J|   SI2| 63.3|  58| 335|4.34|4.35|2.75|
3.1|
| 0.24|Very Good|    J|  VVS2| 62.8|  57| 336|3.94|3.96|2.48|
2.4|
| 0.24|Very Good|    I|  VVS1| 62.3|  57| 336|3.95|3.98|2.47|
2.4|
| 0.26|Very Good|    H|   SI1| 61.9|  55| 337|4.07|4.11|2.53|
2.6|
| 0.22|    Fair|    E|   VS2| 65.1|  61| 337|3.87|3.78|2.49|
2.2|
| 0.23|Very Good|    H|   VS1| 59.4|  61| 338|
4|4.05|2.39|2.3000000000000003|
| 0.3|    Good|    J|   SI1|  64|  55| 339|4.25|4.28|2.73|
3.0|
| 0.23|    Ideal|    J|   VS1| 62.8|  56| 340|3.93|
3.9|2.46|2.3000000000000003|
| 0.22|  Premium|    F|   SI1| 60.4|  61| 342|3.88|3.84|2.33|
2.2|
| 0.31|    Ideal|    J|   SI2| 62.2|  54| 344|4.35|4.37|2.71|
```

```

3.1|
| 0.2| Premium| E| SI2| 60.2| 62| 345|3.79|3.75|2.27|
2.0|
| 0.32| Premium| E| I1| 60.9| 58| 345|4.38|4.42|2.68|
3.2|
| 0.3| Ideal| I| SI2| 62| 54| 348|4.31|4.34|2.68|
3.0|
| 0.3| Good| J| SI1| 63.4| 54| 351|4.23|4.29| 2.7|
3.0|
| 0.3| Good| J| SI1| 63.8| 56| 351|4.23|4.26|2.71|
3.0|
| 0.3|Very Good| J| SI1| 62.7| 59| 351|4.21|4.27|2.66|
3.0|
+-----+-----+-----+-----+-----+-----+-----+-----+-----+
--+
only showing top 20 rows

```

```
[55]: df.show()
```

```

+-----+-----+-----+-----+-----+-----+-----+-----+-----+
|carat|      cut|color|clarity|depth|table|price|  x|  y|  z|
+-----+-----+-----+-----+-----+-----+-----+-----+-----+
| 0.23|   Ideal| E| SI2| 61.5| 55| 326|3.95|3.98|2.43|
| 0.21| Premium| E| SI1| 59.8| 61| 326|3.89|3.84|2.31|
| 0.23|   Good| E| VS1| 56.9| 65| 327|4.05|4.07|2.31|
| 0.29| Premium| I| VS2| 62.4| 58| 334| 4.2|4.23|2.63|
| 0.31|   Good| J| SI2| 63.3| 58| 335|4.34|4.35|2.75|
| 0.24|Very Good| J| VVS2| 62.8| 57| 336|3.94|3.96|2.48|
| 0.24|Very Good| I| VVS1| 62.3| 57| 336|3.95|3.98|2.47|
| 0.26|Very Good| H| SI1| 61.9| 55| 337|4.07|4.11|2.53|
| 0.22|   Fair| E| VS2| 65.1| 61| 337|3.87|3.78|2.49|
| 0.23|Very Good| H| VS1| 59.4| 61| 338| 4|4.05|2.39|
| 0.3|   Good| J| SI1| 64| 55| 339|4.25|4.28|2.73|
| 0.23|   Ideal| J| VS1| 62.8| 56| 340|3.93| 3.9|2.46|
| 0.22| Premium| F| SI1| 60.4| 61| 342|3.88|3.84|2.33|
| 0.31|   Ideal| J| SI2| 62.2| 54| 344|4.35|4.37|2.71|
| 0.2| Premium| E| SI2| 60.2| 62| 345|3.79|3.75|2.27|
| 0.32| Premium| E| I1| 60.9| 58| 345|4.38|4.42|2.68|
| 0.3|   Ideal| I| SI2| 62| 54| 348|4.31|4.34|2.68|
| 0.3|   Good| J| SI1| 63.4| 54| 351|4.23|4.29| 2.7|
| 0.3|   Good| J| SI1| 63.8| 56| 351|4.23|4.26|2.71|
| 0.3|Very Good| J| SI1| 62.7| 59| 351|4.21|4.27|2.66|
+-----+-----+-----+-----+-----+-----+-----+-----+
only showing top 20 rows

```

```
[63]: df2 = df.withColumn("carat10x",df['carat'] * 10)
```

```
[65]: df2.show()
```

```
+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
--+
|carat|      cut|color|clarity|depth|table|price|  x|  y|  z|
carat10x|
+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
--+
| 0.23|   Ideal|   E|   SI2| 61.5|  55|    326|3.95|3.98|2.43|2.3000000000000003|
326|3.95|3.98|2.43|2.3000000000000003|
| 0.21|  Premium|   E|   SI1| 59.8|  61|    326|3.89|3.84|2.31|
2.1|
| 0.23|    Good|   E|   VS1| 56.9|  65|    327|4.05|4.07|2.31|2.3000000000000003|
327|4.05|4.07|2.31|2.3000000000000003|
| 0.29|  Premium|   I|   VS2| 62.4|  58|    334| 4.2|4.23|2.63|
2.9|
| 0.31|    Good|   J|   SI2| 63.3|  58|    335|4.34|4.35|2.75|
3.1|
| 0.24|Very Good|   J|  VVS2| 62.8|  57|    336|3.94|3.96|2.48|
2.4|
| 0.24|Very Good|   I|  VVS1| 62.3|  57|    336|3.95|3.98|2.47|
2.4|
| 0.26|Very Good|   H|   SI1| 61.9|  55|    337|4.07|4.11|2.53|
2.6|
| 0.22|    Fair|   E|   VS2| 65.1|  61|    337|3.87|3.78|2.49|
2.2|
| 0.23|Very Good|   H|   VS1| 59.4|  61|    338|
4|4.05|2.39|2.3000000000000003|
| 0.3|    Good|   J|   SI1|  64|  55|    339|4.25|4.28|2.73|
3.0|
| 0.23|   Ideal|   J|   VS1| 62.8|  56|    340|3.93|
3.9|2.46|2.3000000000000003|
| 0.22|  Premium|   F|   SI1| 60.4|  61|    342|3.88|3.84|2.33|
2.2|
| 0.31|   Ideal|   J|   SI2| 62.2|  54|    344|4.35|4.37|2.71|
3.1|
| 0.2|  Premium|   E|   SI2| 60.2|  62|    345|3.79|3.75|2.27|
2.0|
| 0.32|  Premium|   E|    I1| 60.9|  58|    345|4.38|4.42|2.68|
3.2|
| 0.3|   Ideal|   I|   SI2|  62|  54|    348|4.31|4.34|2.68|
3.0|
| 0.3|    Good|   J|   SI1| 63.4|  54|    351|4.23|4.29| 2.7|
3.0|
| 0.3|    Good|   J|   SI1| 63.8|  56|    351|4.23|4.26|2.71|
```

```

3.0|
| 0.3|Very Good|    J|    SI1| 62.7|    59|   351|4.21|4.27|2.66|
3.0|
+-----+-----+-----+-----+-----+-----+-----+-----+-----+
--+
only showing top 20 rows

```

```

[67]: # Delete/Drop A Column
df2.drop('carat10x').show()

```

```

+-----+-----+-----+-----+-----+-----+-----+-----+-----+
|carat|      cut|color|clarity|depth|table|price|  x|  y|  z|
+-----+-----+-----+-----+-----+-----+-----+-----+-----+
| 0.23|   Ideal|  E|   SI2| 61.5|   55|  326|3.95|3.98|2.43|
| 0.21|  Premium|  E|   SI1| 59.8|   61|  326|3.89|3.84|2.31|
| 0.23|    Good|  E|   VS1| 56.9|   65|  327|4.05|4.07|2.31|
| 0.29|  Premium|  I|   VS2| 62.4|   58|  334| 4.2|4.23|2.63|
| 0.31|    Good|  J|   SI2| 63.3|   58|  335|4.34|4.35|2.75|
| 0.24|Very Good|  J|  VVS2| 62.8|   57|  336|3.94|3.96|2.48|
| 0.24|Very Good|  I|  VVS1| 62.3|   57|  336|3.95|3.98|2.47|
| 0.26|Very Good|  H|   SI1| 61.9|   55|  337|4.07|4.11|2.53|
| 0.22|    Fair|  E|   VS2| 65.1|   61|  337|3.87|3.78|2.49|
| 0.23|Very Good|  H|   VS1| 59.4|   61|  338|  4|4.05|2.39|
| 0.3|    Good|  J|   SI1|  64|   55|  339|4.25|4.28|2.73|
| 0.23|   Ideal|  J|   VS1| 62.8|   56|  340|3.93| 3.9|2.46|
| 0.22|  Premium|  F|   SI1| 60.4|   61|  342|3.88|3.84|2.33|
| 0.31|   Ideal|  J|   SI2| 62.2|   54|  344|4.35|4.37|2.71|
| 0.2|  Premium|  E|   SI2| 60.2|   62|  345|3.79|3.75|2.27|
| 0.32|  Premium|  E|    I1| 60.9|   58|  345|4.38|4.42|2.68|
| 0.3|   Ideal|  I|   SI2|  62|   54|  348|4.31|4.34|2.68|
| 0.3|    Good|  J|   SI1| 63.4|   54|  351|4.23|4.29| 2.7|
| 0.3|    Good|  J|   SI1| 63.8|   56|  351|4.23|4.26|2.71|
| 0.3|Very Good|  J|   SI1| 62.7|   59|  351|4.21|4.27|2.66|
+-----+-----+-----+-----+-----+-----+-----+-----+
only showing top 20 rows

```

0.0.4 GroupBy

- value counts
- aggregate

```

[69]: # Value Counts
# df['cut'].value_counts()
# df.groupby('cut')[].size()
df.groupby('cut').count().show()

```



```

+-----+-----+
|      cut|count|
+-----+-----+
|  Premium|13791|
|    Ideal|21551|
|    Good| 4906|
|    Fair| 1610|
|Very Good|12082|
+-----+-----+

```

```

[70]: # More Groupby
df.groupby('price').mean().show()

```

```

+-----+
|price|
+-----+
| 2904|
| 3210|
| 3414|
| 3606|
| 3959|
| 4032|
| 4821|
| 4937|
| 5325|
| 5645|
| 5925|
| 6194|
| 6240|
| 6613|
| 6731|
| 7273|
| 7711|
| 7762|
| 9009|
| 9030|
+-----+

```

only showing top 20 rows

```

[71]: # Sum of A groupby
df.groupby('price').sum().show()

```

```

+-----+
|price|
+-----+
| 2904|

```

```
| 3210|
| 3414|
| 3606|
| 3959|
| 4032|
| 4821|
| 4937|
| 5325|
| 5645|
| 5925|
| 6194|
| 6240|
| 6613|
| 6731|
| 7273|
| 7711|
| 7762|
| 9009|
| 9030|
```

```
+-----+
```

only showing top 20 rows

```
[ ]: # # Aggregation
      # df.groupby('carat').agg('col':'sum')
```

```
[73]: df.show()
```

```
+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
|carat|      cut|color|clarity|depth|table|price|  x|  y|  z|
+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
| 0.23|   Ideal|  E|   SI2| 61.5|  55|  326|3.95|3.98|2.43|
| 0.21|  Premium|  E|   SI1| 59.8|  61|  326|3.89|3.84|2.31|
| 0.23|    Good|  E|   VS1| 56.9|  65|  327|4.05|4.07|2.31|
| 0.29|  Premium|  I|   VS2| 62.4|  58|  334| 4.2|4.23|2.63|
| 0.31|    Good|  J|   SI2| 63.3|  58|  335|4.34|4.35|2.75|
| 0.24|Very Good|  J|  VVS2| 62.8|  57|  336|3.94|3.96|2.48|
| 0.24|Very Good|  I|  VVS1| 62.3|  57|  336|3.95|3.98|2.47|
| 0.26|Very Good|  H|   SI1| 61.9|  55|  337|4.07|4.11|2.53|
| 0.22|    Fair|  E|   VS2| 65.1|  61|  337|3.87|3.78|2.49|
| 0.23|Very Good|  H|   VS1| 59.4|  61|  338|  4|4.05|2.39|
|  0.3|    Good|  J|   SI1|  64|  55|  339|4.25|4.28|2.73|
| 0.23|   Ideal|  J|   VS1| 62.8|  56|  340|3.93| 3.9|2.46|
| 0.22|  Premium|  F|   SI1| 60.4|  61|  342|3.88|3.84|2.33|
| 0.31|   Ideal|  J|   SI2| 62.2|  54|  344|4.35|4.37|2.71|
|  0.2|  Premium|  E|   SI2| 60.2|  62|  345|3.79|3.75|2.27|
| 0.32|  Premium|  E|    I1| 60.9|  58|  345|4.38|4.42|2.68|
|  0.3|   Ideal|  I|   SI2|  62|  54|  348|4.31|4.34|2.68|
```

```
| 0.3|      Good|    J|    SI1| 63.4|    54|  351|4.23|4.29| 2.7|
| 0.3|      Good|    J|    SI1| 63.8|    56|  351|4.23|4.26|2.71|
| 0.3|Very Good|    J|    SI1| 62.7|    59|  351|4.21|4.27|2.66|
+-----+-----+-----+-----+-----+-----+-----+-----+
only showing top 20 rows
```

```
[74]: df.columns
```

```
[74]: ['carat', 'cut', 'color', 'clarity', 'depth', 'table', 'price', 'x', 'y', 'z']
```

```
[75]: # Rearrange Columns
df.select('carat', 'color', 'clarity', 'depth', 'table', 'price', 'x', 'y', 'z',
         ↪ 'z', 'cut').show()
```

```
+-----+-----+-----+-----+-----+-----+-----+-----+
|carat|color|clarity|depth|table|price|  x|  y|  z|      cut|
+-----+-----+-----+-----+-----+-----+-----+-----+
| 0.23|    E|    SI2| 61.5|  55|  326|3.95|3.98|2.43|    Ideal|
| 0.21|    E|    SI1| 59.8|  61|  326|3.89|3.84|2.31|  Premium|
| 0.23|    E|    VS1| 56.9|  65|  327|4.05|4.07|2.31|    Good|
| 0.29|    I|    VS2| 62.4|  58|  334| 4.2|4.23|2.63|  Premium|
| 0.31|    J|    SI2| 63.3|  58|  335|4.34|4.35|2.75|    Good|
| 0.24|    J|   VVS2| 62.8|  57|  336|3.94|3.96|2.48|Very Good|
| 0.24|    I|   VVS1| 62.3|  57|  336|3.95|3.98|2.47|Very Good|
| 0.26|    H|    SI1| 61.9|  55|  337|4.07|4.11|2.53|Very Good|
| 0.22|    E|    VS2| 65.1|  61|  337|3.87|3.78|2.49|    Fair|
| 0.23|    H|    VS1| 59.4|  61|  338|  4|4.05|2.39|Very Good|
| 0.3|    J|    SI1|  64|  55|  339|4.25|4.28|2.73|    Good|
| 0.23|    J|    VS1| 62.8|  56|  340|3.93| 3.9|2.46|    Ideal|
| 0.22|    F|    SI1| 60.4|  61|  342|3.88|3.84|2.33|  Premium|
| 0.31|    J|    SI2| 62.2|  54|  344|4.35|4.37|2.71|    Ideal|
| 0.2|    E|    SI2| 60.2|  62|  345|3.79|3.75|2.27|  Premium|
| 0.32|    E|    I1| 60.9|  58|  345|4.38|4.42|2.68|  Premium|
| 0.3|    I|    SI2|  62|  54|  348|4.31|4.34|2.68|    Ideal|
| 0.3|    J|    SI1| 63.4|  54|  351|4.23|4.29| 2.7|    Good|
| 0.3|    J|    SI1| 63.8|  56|  351|4.23|4.26|2.71|    Good|
| 0.3|    J|    SI1| 62.7|  59|  351|4.21|4.27|2.66|Very Good|
+-----+-----+-----+-----+-----+-----+-----+-----+
only showing top 20 rows
```

```
[76]: # Assign DF to a New DF
new_df = df.select('carat', 'color', 'clarity', 'depth', 'table', 'price', 'x', 'y',
         ↪ 'y', 'z', 'cut')
```

```
[77]: new_df
```

```
[77]: DataFrame[carat: string, color: string, clarity: string, depth: string, table: string, price: string, x: string, y: string, z: string, cut: string]
```

```
[78]: # Check Datatype
new_df.dtypes
```

```
[78]: [('carat', 'string'),
      ('color', 'string'),
      ('clarity', 'string'),
      ('depth', 'string'),
      ('table', 'string'),
      ('price', 'string'),
      ('x', 'string'),
      ('y', 'string'),
      ('z', 'string'),
      ('cut', 'string')]
```

```
[80]: # Check For the Schema
df.printSchema()
```

```
root
|-- carat: string (nullable = true)
|-- cut: string (nullable = true)
|-- color: string (nullable = true)
|-- clarity: string (nullable = true)
|-- depth: string (nullable = true)
|-- table: string (nullable = true)
|-- price: string (nullable = true)
|-- x: string (nullable = true)
|-- y: string (nullable = true)
|-- z: string (nullable = true)
```

```
[81]: # Check type of DF
type(df)
```

```
[81]: pyspark.sql.dataframe.DataFrame
```

0.0.5 Saving DataFrames as CSV,parquet etc

```
[83]: # Save
new_df.write.format('csv').option('header','true').save("diamond_clean.csv")
```

```
[84]: !ls
```

```
diamond_clean.csv  diamonds.csv  'PySpark Crash Course.ipynb'  work
```

```
[ ]: # Save as parquet
new_df.write.format('parquet').save("diamond_clean.parquet")
```

0.0.6 Making SQL Queries

- parse in the spark.SparkContext
- sqlContext

```
[85]: from pyspark.sql import SQLContext
```

```
[92]: dir(spark)
```

```
[92]: ['Builder',
      '__class__',
      '__delattr__',
      '__dict__',
      '__dir__',
      '__doc__',
      '__enter__',
      '__eq__',
      '__exit__',
      '__format__',
      '__ge__',
      '__getattribute__',
      '__gt__',
      '__hash__',
      '__init__',
      '__init_subclass__',
      '__le__',
      '__lt__',
      '__module__',
      '__ne__',
      '__new__',
      '__reduce__',
      '__reduce_ex__',
      '__repr__',
      '__setattr__',
      '__sizeof__',
      '__str__',
      '__subclasshook__',
      '__weakref__',
      '_activeSession',
      '_convert_from_pandas',
      '_createFromLocal',
      '_createFromRDD',
      '_create_dataframe',
      '_create_from_pandas_with_arrow',
```

```

'_create_shell_session',
'_get_numpy_record_dtype',
'_inferSchema',
'_inferSchemaFromList',
'_instantiatedSession',
'_jsc',
'_jsparkSession',
'_jvm',
'_jwrapped',
'_repr_html_',
'_sc',
'_wrapped',
'builder',
'catalog',
'conf',
'createDataFrame',
'getActiveSession',
'newSession',
'range',
'read',
'readStream',
'sparkContext',
'sql',
'stop',
'streams',
'table',
'udf',
'version']

```

```
[93]: # Create A Spark Context From the Spark Session
sc = spark.sparkContext
```

```
[94]: # Parse into the SQLContext
sqlContext = SQLContext(sc)
```

```
[96]: # Register Current DataFrame As Temporal Table
df.registerTempTable("DiamondsTable")
```

```
[97]: # Making QUeries
sqlContext.sql('SELECT * FROM DiamondsTable').show()
```

```

+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
|carat|      cut|color|clarity|depth|table|price|  x|  y|  z|
+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
| 0.23|   Ideal|   E|   SI2| 61.5|  55|  326|3.95|3.98|2.43|
| 0.21| Premium|   E|   SI1| 59.8|  61|  326|3.89|3.84|2.31|
| 0.23|    Good|   E|   VS1| 56.9|  65|  327|4.05|4.07|2.31|

```

0.29	Premium	I	VS2	62.4	58	334	4.2	4.23	2.63
0.31	Good	J	SI2	63.3	58	335	4.34	4.35	2.75
0.24	Very Good	J	VVS2	62.8	57	336	3.94	3.96	2.48
0.24	Very Good	I	VVS1	62.3	57	336	3.95	3.98	2.47
0.26	Very Good	H	SI1	61.9	55	337	4.07	4.11	2.53
0.22	Fair	E	VS2	65.1	61	337	3.87	3.78	2.49
0.23	Very Good	H	VS1	59.4	61	338	4	4.05	2.39
0.3	Good	J	SI1	64	55	339	4.25	4.28	2.73
0.23	Ideal	J	VS1	62.8	56	340	3.93	3.9	2.46
0.22	Premium	F	SI1	60.4	61	342	3.88	3.84	2.33
0.31	Ideal	J	SI2	62.2	54	344	4.35	4.37	2.71
0.2	Premium	E	SI2	60.2	62	345	3.79	3.75	2.27
0.32	Premium	E	I1	60.9	58	345	4.38	4.42	2.68
0.3	Ideal	I	SI2	62	54	348	4.31	4.34	2.68
0.3	Good	J	SI1	63.4	54	351	4.23	4.29	2.7
0.3	Good	J	SI1	63.8	56	351	4.23	4.26	2.71
0.3	Very Good	J	SI1	62.7	59	351	4.21	4.27	2.66

+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+

only showing top 20 rows

```
[98]: # Can also use it to work with DataFrames
      dir(sqlContext)
```

```
[98]: ['__class__',
      '__delattr__',
      '__dict__',
      '__dir__',
      '__doc__',
      '__eq__',
      '__format__',
      '__ge__',
      '__getattr__',
      '__gt__',
      '__hash__',
      '__init__',
      '__init_subclass__',
      '__le__',
      '__lt__',
      '__module__',
      '__ne__',
      '__new__',
      '__reduce__',
      '__reduce_ex__',
      '__repr__',
      '__setattr__',
      '__sizeof__']
```

```

'__str__',
'__subclasshook__',
'__weakref__',
'_conf',
'_inferSchema',
'_instantiatedContext',
'_jsc',
'_jdbcContext',
'_jvm',
'_sc',
'_ssql_ctx',
'cacheTable',
'clearCache',
'createDataFrame',
'createExternalTable',
'dropTempTable',
'getConf',
'getOrCreate',
'newSession',
'range',
'read',
'readStream',
'registerDataFrameAsTable',
'registerFunction',
'registerJavaFunction',
'setConf',
'sparkSession',
'sql',
'streams',
'table',
'tableNames',
'tables',
'udf',
'uncacheTable']

```

```

[99]: # Thanks For Watching
      # Jesus Saves @JCharisTech
      # Jesse E.Agbe(JCharis)
      # 2021

```

```
[ ]:
```